

In [0]:

```
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
import matplotlib.pyplot as plt
import numpy as np
import time
```

In [0]:

```
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [0]:

```
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>  
11493376/11490434 [=====] - 1s 0us/step

In [0]:

```
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [0]:

```
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  249 253 249  64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0]
```

In [0]:

In [0]:

[illegible]

[illegible]

In [0]:

```
print("Class label of first image :", y_train[0])
```

Class label of first image : 5

In [0]:

```
from keras.utils import np_utils
Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)
print("After converting the output into a vector : ",Y_train[0])
```

After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

In [0]:

```
from keras.models import Sequential
from keras.layers import Dense, Activation
```

In [0]:

```
output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

For (784 - 488 - 256 - 10) Architecture

In [0]:

```
from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization

model = Sequential()

model.add(Dense(488, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(output_dim, activation='softmax'))

model.summary()
```

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 488)	383080
batch_normalization_5 (Batch Normalization)	(None, 488)	1952
dropout_5 (Dropout)	(None, 488)	0
dense_8 (Dense)	(None, 256)	125184
batch_normalization_6 (Batch Normalization)	(None, 256)	1024
dropout_6 (Dropout)	(None, 256)	0
dense_9 (Dense)	(None, 10)	2570

```
=====
Total params: 513,810
Trainable params: 512,322
Non-trainable params: 1,488
=====
```

In [0]:

```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 11s 179us/step - loss: 0.4587 - acc: 0.8601 - val_loss: 0.1523 - val_acc: 0.9538
Epoch 2/20
60000/60000 [=====] - 10s 163us/step - loss: 0.2336 - acc: 0.9285 - val_loss: 0.1202 - val_acc: 0.9620
Epoch 3/20
60000/60000 [=====] - 10s 164us/step - loss: 0.1864 - acc: 0.9437 - val_loss: 0.1012 - val_acc: 0.9686
Epoch 4/20
60000/60000 [=====] - 10s 166us/step - loss: 0.1616 - acc: 0.9516 - val_loss: 0.0855 - val_acc: 0.9728
Epoch 5/20
60000/60000 [=====] - 10s 165us/step - loss: 0.1467 - acc: 0.9553 - val_loss: 0.0863 - val_acc: 0.9723
Epoch 6/20
60000/60000 [=====] - 10s 164us/step - loss: 0.1282 - acc: 0.9610 - val_loss: 0.0823 - val_acc: 0.9739
Epoch 7/20
60000/60000 [=====] - 10s 165us/step - loss: 0.1258 - acc: 0.9615 - val_loss: 0.0802 - val_acc: 0.9740
Epoch 8/20
60000/60000 [=====] - 10s 165us/step - loss: 0.1143 - acc: 0.9644 - val_loss: 0.0752 - val_acc: 0.9764
Epoch 9/20
60000/60000 [=====] - 10s 165us/step - loss: 0.1080 - acc: 0.9660 - val_loss: 0.0709 - val_acc: 0.9780
Epoch 10/20
60000/60000 [=====] - 10s 167us/step - loss: 0.1030 - acc: 0.9677 - val_loss: 0.0679 - val_acc: 0.9783
Epoch 11/20
60000/60000 [=====] - 10s 167us/step - loss: 0.0971 - acc: 0.9699 - val_loss: 0.0657 - val_acc: 0.9796
Epoch 12/20
60000/60000 [=====] - 10s 168us/step - loss: 0.0935 - acc: 0.9713 - val_loss: 0.0674 - val_acc: 0.9804
Epoch 13/20
60000/60000 [=====] - 10s 166us/step - loss: 0.0841 - acc: 0.9735 - val_loss: 0.0637 - val_acc: 0.9802
Epoch 14/20
60000/60000 [=====] - 10s 166us/step - loss: 0.0854 - acc: 0.9733 - val_loss: 0.0637 - val_acc: 0.9807
Epoch 15/20
60000/60000 [=====] - 10s 168us/step - loss: 0.0810 - acc: 0.9746 - val_loss: 0.0652 - val_acc: 0.9819
Epoch 16/20
60000/60000 [=====] - 10s 168us/step - loss: 0.0779 - acc: 0.9753 - val_loss: 0.0603 - val_acc: 0.9826
Epoch 17/20
60000/60000 [=====] - 10s 166us/step - loss: 0.0762 - acc: 0.9757 - val_loss: 0.0608 - val_acc: 0.9825
Epoch 18/20
60000/60000 [=====] - 10s 168us/step - loss: 0.0709 - acc: 0.9774 - val_loss: 0.0603 - val_acc: 0.9830
Epoch 19/20
60000/60000 [=====] - 10s 166us/step - loss: 0.0658 - acc: 0.9783 - val_loss: 0.0621 - val_acc: 0.9833
Epoch 20/20
60000/60000 [=====] - 10s 166us/step - loss: 0.0681 - acc: 0.9785 - val_loss: 0.0618 - val_acc: 0.9819
```

In [0]:

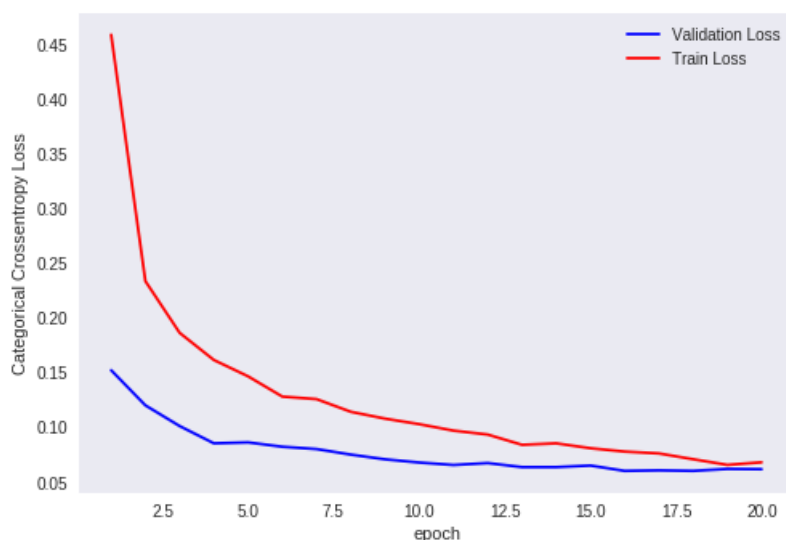
```
score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.0618182118332712

Test accuracy: 0.9819



For (784 - 515 - 319 - 127 - 10) 3 hidden layer Architecture

In [0]:

```
from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization

model = Sequential()

model.add(Dense(515, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(319, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(127, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(output_dim, activation='softmax'))

model.summary()
```

Layer (type)	Output Shape	Param #
dense_16 (Dense)	(None, 515)	404275

batch_normalization_11 (Batch Normalization)	(None, 515)	2060
dropout_11 (Dropout)	(None, 515)	0
dense_17 (Dense)	(None, 319)	164604
batch_normalization_12 (Batch Normalization)	(None, 319)	1276
dropout_12 (Dropout)	(None, 319)	0
dense_18 (Dense)	(None, 127)	40640
batch_normalization_13 (Batch Normalization)	(None, 127)	508
dropout_13 (Dropout)	(None, 127)	0
dense_19 (Dense)	(None, 10)	1280
=====		
Total params: 614,643		
Trainable params: 612,721		
Non-trainable params: 1,922		

```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```

oss: 0.0592 - val_acc: 0.9819
Epoch 16/20
60000/60000 [=====] - 13s 224us/step - loss: 0.0746 - acc: 0.9772 - val_l
oss: 0.0580 - val_acc: 0.9830
Epoch 17/20
60000/60000 [=====] - 13s 218us/step - loss: 0.0746 - acc: 0.9775 - val_l
oss: 0.0608 - val_acc: 0.9827
Epoch 18/20
60000/60000 [=====] - 13s 215us/step - loss: 0.0704 - acc: 0.9783 - val_l
oss: 0.0616 - val_acc: 0.9830
Epoch 19/20
60000/60000 [=====] - 13s 215us/step - loss: 0.0646 - acc: 0.9798 - val_l
oss: 0.0595 - val_acc: 0.9841
Epoch 20/20
60000/60000 [=====] - 13s 215us/step - loss: 0.0648 - acc: 0.9804 - val_l
oss: 0.0529 - val_acc: 0.9858

```

In [0]:

```


score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.052927311327529605  
Test accuracy: 0.9858



For (784 - 600 - 567-428-385 -250- 10) 5 hidden layer Architecture

In [0]:

```

from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization

model = Sequential()

model.add(Dense(600, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(m
ean=0.0, stddev=0.039, seed=None)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(567, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(m
ean=0.0, stddev=0.039, seed=None)))
model.add(BatchNormalization())

```



```

model.add(Dropout(0.5))

model.add(Dense(428, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(385, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(250, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(output_dim, activation='softmax'))

model.summary()

```

Layer (type)	Output Shape	Param #
dense_20 (Dense)	(None, 600)	471000
batch_normalization_14 (Batch Normalization)	(None, 600)	2400
dropout_14 (Dropout)	(None, 600)	0
dense_21 (Dense)	(None, 567)	340767
batch_normalization_15 (Batch Normalization)	(None, 567)	2268
dropout_15 (Dropout)	(None, 567)	0
dense_22 (Dense)	(None, 428)	243104
batch_normalization_16 (Batch Normalization)	(None, 428)	1712
dropout_16 (Dropout)	(None, 428)	0
dense_23 (Dense)	(None, 385)	165165
batch_normalization_17 (Batch Normalization)	(None, 385)	1540
dropout_17 (Dropout)	(None, 385)	0
dense_24 (Dense)	(None, 250)	96500
batch_normalization_18 (Batch Normalization)	(None, 250)	1000
dropout_18 (Dropout)	(None, 250)	0
dense_25 (Dense)	(None, 10)	2510
Total params: 1,327,966		
Trainable params: 1,323,506		
Non-trainable params: 4,460		

In [0]:

```

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 29s 481us/step - loss: 0.0717 - acc: 0.9793 - val_loss: 0.0594 - val_acc: 0.9850
Epoch 2/20
60000/60000 [=====] - 26s 436us/step - loss: 0.0706 - acc: 0.9798 - val_loss: 0.0594 - val_acc: 0.9850

```

```

60000/60000 [-----] - 26s 435us/step - loss: 0.0695 - acc: 0.9798 - val_loss: 0.0578 - val_acc: 0.9838
Epoch 3/20
60000/60000 [=====] - 26s 435us/step - loss: 0.0695 - acc: 0.9798 - val_loss: 0.0578 - val_acc: 0.9838
Epoch 4/20
60000/60000 [=====] - 26s 433us/step - loss: 0.0631 - acc: 0.9816 - val_loss: 0.0557 - val_acc: 0.9848
Epoch 5/20
60000/60000 [=====] - 26s 433us/step - loss: 0.0616 - acc: 0.9816 - val_loss: 0.0511 - val_acc: 0.9857
Epoch 6/20
60000/60000 [=====] - 26s 434us/step - loss: 0.0603 - acc: 0.9823 - val_loss: 0.0531 - val_acc: 0.9849
Epoch 7/20
60000/60000 [=====] - 26s 435us/step - loss: 0.0605 - acc: 0.9827 - val_loss: 0.0578 - val_acc: 0.9846
Epoch 8/20
60000/60000 [=====] - 26s 438us/step - loss: 0.0570 - acc: 0.9829 - val_loss: 0.0587 - val_acc: 0.9849
Epoch 9/20
60000/60000 [=====] - 26s 436us/step - loss: 0.0571 - acc: 0.9836 - val_loss: 0.0520 - val_acc: 0.9857
Epoch 10/20
60000/60000 [=====] - 27s 450us/step - loss: 0.0540 - acc: 0.9841 - val_loss: 0.0541 - val_acc: 0.9854
Epoch 11/20
60000/60000 [=====] - 26s 437us/step - loss: 0.0513 - acc: 0.9839 - val_loss: 0.0558 - val_acc: 0.9854
Epoch 12/20
60000/60000 [=====] - 26s 440us/step - loss: 0.0545 - acc: 0.9838 - val_loss: 0.0535 - val_acc: 0.9853
Epoch 13/20
60000/60000 [=====] - 26s 436us/step - loss: 0.0501 - acc: 0.9851 - val_loss: 0.0533 - val_acc: 0.9864
Epoch 14/20
60000/60000 [=====] - 26s 437us/step - loss: 0.0512 - acc: 0.9850 - val_loss: 0.0505 - val_acc: 0.9866
Epoch 15/20
60000/60000 [=====] - 26s 439us/step - loss: 0.0490 - acc: 0.9854 - val_loss: 0.0530 - val_acc: 0.9863
Epoch 16/20
60000/60000 [=====] - 26s 439us/step - loss: 0.0438 - acc: 0.9866 - val_loss: 0.0533 - val_acc: 0.9859
Epoch 17/20
60000/60000 [=====] - 26s 438us/step - loss: 0.0455 - acc: 0.9862 - val_loss: 0.0528 - val_acc: 0.9857
Epoch 18/20
60000/60000 [=====] - 26s 436us/step - loss: 0.0439 - acc: 0.9868 - val_loss: 0.0504 - val_acc: 0.9869
Epoch 19/20
60000/60000 [=====] - 27s 442us/step - loss: 0.0427 - acc: 0.9874 - val_loss: 0.0546 - val_acc: 0.9857
Epoch 20/20
60000/60000 [=====] - 27s 442us/step - loss: 0.0432 - acc: 0.9871 - val_loss: 0.0531 - val_acc: 0.9861

```

In [0]:

```


score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.05311666791146854  
Test accuracy: 0.9861



In [1]:

```
from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["Model", "AUC"]

x.add_row(["(784 - 488 - 256 - 10) Architecture", "98.19%"])
x.add_row(["(784 - 515 - 319 - 127 - 10) Architecture", "98.58%"])
x.add_row(["(784 - 600 - 567-428-385 -250- 10) ", "98.61%"])

print(x)
```

Model	AUC
(784 - 488 - 256 - 10) Architecture	98.19%
(784 - 515 - 319 - 127 - 10) Architecture	98.58%
(784 - 600 - 567-428-385 -250- 10)	98.61%

Observation: 1) Accuracy increasing slightly as number of hidden layer increases. 2) As in graph of third architecture of number of 5 hidden layers its overfitting after 13 epochs.